

- When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e.,  $n < 30$ ). It turns out, that how “far off” the  $t$ -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

(a) Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6}}_{\text{error}} f_Z(t),$$

where  $f_Z(\cdot)$  and  $F_Z(\cdot)$  are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the  $p$ -value when testing  $H_0 : \mu_X = 0; H_a : \mu_X < 0$  using the zebra finch further data?

**Answer:** The error is approximately  $-1.225 \times 10^{-13}$


```
#####
further.dat <- finches.dat$further
mu.0 <- 0
t.further <- t.test(x = further.dat,
                    alternative = "less",
                    mu = mu.0)

# t-value
t <- t.further$statistic
n <- length(further.dat)
gauss.pdf <- dnorm(x = t)
gauss.cdf <- pnorm(q = t)
(error <- ((skewness(further.dat) / sqrt(n)) * ((2*t^2 + 1)/6)) * gauss.pdf)

##          t
## -1.226006e-13
#####
```

- Compute the error for  $t$  statistics from -10 to 10 and plot a line that shows the error across  $t$ . Continue to use the skewness and the sample size for the zebra finch further data.

```
#####
tVals <- seq(-10, 10, 0.1)
t.errors = tibble(t = numeric(length(tVals)), error = numeric(length(tVals)))
i = 1
for(t in tVals){
  t.errors$t[i] <- t
  gauss.pdf <- dnorm(x = t)
  gauss.cdf <- pnorm(q = t)
  curr.error <- ((skewness(further.dat) / sqrt(n)) * ((2*t^2 + 1)/6)) * gauss.pdf
  t.errors$error[i] <- curr.error
  i <- i+1
}
first.p <- ggplot(x = t.errors) +
  geom_line(aes(x = t.errors$t, y = t.errors$error)) +
  theme_bw() +
  labs(y = "Error",
       x = "t-value",
       title = "Edgeworth Approximation Error")
#####
```



- (c) Suppose we wanted to have a tail probability within 10% of the desired  $\alpha = 0.05$ . Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{set}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left( \frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

**Answer:** We need a sample size of approximately 520.

```
#####
t <- qnorm(0.05)
n <- length(further.dat)
gauss.pdf <- dnorm(x = t)
gauss.cdf <- pnorm(q = t)
alpha <- 0.05
skew <- skewness(further.dat)
(size <- (((skew/(6*(0.1*alpha))) * ((2*t^2) + 1) * gauss.pdf)^2))

## [1] 520.8876
#####
```

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the  $T$  statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where  $\bar{x}_r$  is the mean computed on the  $r^{th}$  resample and  $s$  is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the

average (i.e., here ensure the shifted resamples are 0 on average, corresponding to  $t = 0$ , for each case).

```
#####
# re-sampling for further data
n <- length(further.dat)
R <- 10000
resamples.further <- tibble(t.stat = numeric(R), mean = numeric(R))
s <- sd(further.dat)

for(i in 1:R){
  curr.sample <- sample(x = further.dat,
                        size = n,
                        replace = T)
  t.stat <- mean(curr.sample) / (s/sqrt(n))
  resamples.further$t.stat[i] <- t.stat
  resamples.further$mean[i] <- mean(curr.sample)
}

#####
# further bootstrapping
f.delta <- mean(resamples.further$t.stat) - 0
resamples.null.further <- resamples.further |>
mutate(t.shifted = resamples.further$t.stat - f.delta)
# re sampling for closer data
closer.dat <- finches.dat$closer
n <- length(closer.dat)
R <- 10000
s <- sd(closer.dat)
resamples.closer <- tibble(t.stat = numeric(R), mean = numeric(R))

for (i in 1:R){
  curr.sample <- sample(x = closer.dat,
                        size = n,
                        replace = T)
  t.stat <- mean(curr.sample) / (s/sqrt(n))
  resamples.closer$t.stat[i] <- t.stat
  resamples.closer$mean[i] <- mean(curr.sample)
}

#####
# closer bootstrapping
c.delta <- mean(resamples.closer$t.stat) - 0
resamples.null.closer <- resamples.closer |>
mutate(t.shifted = resamples.closer$t.stat - c.delta)
#####
# re sampling for difference data
difference.dat <- finches.dat$diff
n <- length(closer.dat)
R <- 10000
s <- sd(difference.dat)
resamples.difference <- tibble(t.stat = numeric(R), mean = numeric(R))

for (i in 1:R){
  curr.sample <- sample(x = difference.dat,
                        size = n,
                        replace = T)
  t.stat <- mean(curr.sample) / (s/sqrt(n))
  resamples.difference$t.stat[i] <- t.stat
  resamples.difference$mean[i] <- mean(curr.sample)
}

#####
# diff bootstrapping
d.delta <- mean(resamples.difference$t.stat) - 0
resamples.null.diff <- resamples.difference |>
mutate(t.shifted = resamples.difference$t.stat - d.delta)
#####
```

(b) Compute the bootstrap  $p$ -value for each test using the shifted resamples.

```
#####
# further
further.boot.p <- mean(resamples.null.further$t.shifted <= f.delta )
further.t.p <- (t.test(x = further.dat, mu = 0, alternative = "less"))$p.value
# closer
closer.boot.p <- mean(resamples.null.closer$t.shifted >= c.delta)
closer.t.p <- (t.test(x = closer.dat, mu = 0, alternative = "greater"))$p.value
# difference
low <- mean(resamples.null.diff$t.shifted <= -d.delta)
high <- mean(resamples.null.diff$t.shifted >= d.delta)
diff.boot.p <- low + high
diff.t.p <- (t.test(x = difference.dat, mu = 0, alternative = "two.sided"))$p.value
```

```
#####
```

data	t.test.p	boot.p
further	0.00	0.00
closer	0.00	0.00
difference	0.00	0.00

Table 1: T-test and Bootstrap p-values

How do these compare to the  $t$ -test  $p$ -values? **Answer:** The bootstrap  $p$ -values are all zero, which is the same as the  $p$ -values from the  $t$ -test.

- (c) What is the 5<sup>th</sup> percentile of the shifted resamples under the null hypothesis? Note this value approximates  $t_{0.05, n-1}$ . Compare these values in each case.

```
#####
further.boot.ptl <- quantile(resamples.null.further$t.shifted, 0.05)
further.t.ptl <- qt(0.05, df = n - 1)

closer.boot.ptl <- quantile(resamples.null.closer$t.shifted, 0.95)
closer.t.ptl <- qt(0.95, df = n - 1)

diff.boot.ptl <- quantile(resamples.null.diff$t.shifted, 0.05)
diff.t.ptl <- qt(0.05, df = n - 1)
percentile.summary <- tibble(data = c("further", "closer", "difference"),
                             t.test.p = c(further.t.ptl, closer.t.ptl, diff.t.ptl),
                             boot.p = c(further.boot.ptl, closer.boot.ptl, diff.boot.ptl))
#####
```

data	t.test.p	boot.p
further	-1.71	-1.67
closer	-1.71	-1.60
difference	-1.71	-1.56

Table 2: T-test and Bootstrap 5<sup>th</sup> percentiles

**Answer:** We can see that percentiles for the two methods are roughly the same, although the bootstrap values are slightly less extreme and closer to 0.

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the  $t$ -test confidence intervals?

```
#####
# further
further.boot.ci <- quantile(resamples.further$mean, c(0.025, 0.975))
further.t.ci <- t.test(further.dat, mu = 0, conf.level = 0.95,
                      alternative = "two.sided")$conf.int

# closer
closer.boot.ci = quantile(resamples.closer$mean, c(0.025, 0.975))
closer.t.ci <- t.test(closer.dat, mu = 0, conf.level = 0.95,
                     alternative = "two.sided")$conf.int

# difference
diff.boot.ci = quantile(resamples.difference$mean, c(0.025, 0.975))
diff.t.ci <- t.test(difference.dat, mu = 0, conf.level = 0.95,
                   alternative = "two.sided")$conf.int

# further
further.boot.ci # bootstrapping confidence interval

##      2.5%      97.5%
## -0.2552970 -0.1545911

further.t.ci # t.test confidence interval

## [1] -0.2565176 -0.1489313
## attr(,"conf.level")
## [1] 0.95

# closer
closer.boot.ci
```

```
##      2.5%      97.5%
## 0.1215823 0.1928475

closer.t.ci

## [1] 0.1173875 0.1950586
## attr(,"conf.level")
## [1] 0.95

# difference
diff.boot.ci

##      2.5%      97.5%
## 0.2810074 0.4433528

diff.t.ci

## [1] 0.2719028 0.4459921
## attr(,"conf.level")
## [1] 0.95

#####
```

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

```
#####
# randomization for further data
shifted.further <- further.dat - mu.0
# perform (randomization) shuffling
further.rand <- tibble(mean = numeric(R))
s <- sd(further.dat)
n <- length(further.dat)

for(i in 1:R){
  curr.rand <- shifted.further *
    sample(x = c(-1, 1),
           size = length(shifted.further),
           replace = T)
  further.rand$mean[i] <- mean(curr.rand)
}

further.rand <- further.rand |>
  mutate(mean = mean + mu.0)

# randomization for closer data
shifted.closer <- closer.dat - mu.0
# perform (randomization) shuffling
closer.rand <- tibble(mean = numeric(R))
s <- sd(closer.dat)
n <- length(closer.dat)
for(i in 1:R){
  curr.rand <- shifted.closer *
    sample(x = c(-1, 1),
           size = length(shifted.closer),
           replace = T)
  closer.rand$mean[i] <- mean(curr.rand)
}
# shift back
closer.rand <- closer.rand |>
  mutate(mean = mean + mu.0)

# randomization for difference data
shifted.difference <- difference.dat - mu.0
# perform (randomization) shuffling
diff.rand <- tibble(mean = numeric(R))
s <- sd(difference.dat)
n <- length(difference.dat)
for(i in 1:R){
  curr.rand <- shifted.difference *
    sample(x = c(-1, 1),
           size = length(shifted.difference),
           replace = T)
  diff.rand$mean[i] <- mean(curr.rand)
}
# shift back
diff.rand <- diff.rand |>
  mutate(mean = mean + mu.0)
```

- (b) Compute the randomization test  $p$ -value for each test.

```
#####
# further data
delta <- abs(mean(further.dat) - mu.0)
low <- mu.0 - delta # mirror
high <- mu.0 + delta # xbar
further.rand.p <- mean(further.rand$mean <= low) +
  mean(further.rand$mean >= high)
# closer data
delta <- abs(mean(closer.dat) - mu.0)
low <- mu.0 - delta # mirror
high <- mu.0 + delta # xbar
closer.rand.p <- mean(closer.rand$mean <= low) +
  mean(closer.rand$mean >= high)
# difference data
delta <- abs(mean(difference.dat) - mu.0)
low <- mu.0 - delta # mirror
high <- mu.0 + delta # xbar
diff.rand.p <- mean(diff.rand$mean <= low) +
  mean(diff.rand$mean >= high)
random.p.summary <- tibble(data = c("further", "closer", "difference"),
  p_value = c(further.rand.p, closer.rand.p, diff.rand.p))
#####
```

data	p_value
further	0.00
closer	0.00
difference	0.00

- (c) Compute the randomization confidence interval by iterating over values of  $\mu_0$ .  
**Hint:** You can “search” for the lower bound from  $Q_1$  and subtracting by 0.0001, and the upper bound using  $Q_3$  and increasing by 0.0001. You will continue until you find the first value for which the two-sided  $p$ -value is greater than or equal to 0.05.  
**Answer:** The confidence interval for the further data is (-0.203 -0.151), the closer data is (0.118, 1.195) and difference data is (0.276, 0.446)

```
#####
# for further data
R <- 1000
mu0.iterate <- 0.001
starting.point <- mean(further.dat)
mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))
  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- further.dat - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)
    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
  mutate(xbars = xbars + mu.lower) # shifting back
  # p-value
  delta <- abs(mean(further.dat) - mu.lower)
  low <- mu.lower - delta # mirror
  high <- mu.lower + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }else{
    mu.lower <- mu.lower - mu0.iterate
  }
}

mu.upper <- starting.point
```

```

repeat{
  rand <- tibble(xbars = rep(NA, R))
  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- further.dat - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  delta <- abs(mean(further.dat) - mu.upper)
  low <- mu.upper - delta # mirror
  high <- mu.upper + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}
further.rand.ci <- c(mu.lower, mu.upper)

# for closer data
R <- 1000
mu0.iterate <- 0.001
starting.point <- mean(closer.dat)
mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))
  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- closer.dat - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back

  # p-value
  delta <- abs(mean(closer.dat) - mu.lower)
  low <- mu.lower - delta # mirror
  high <- mu.lower + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)
  if(p.val < 0.05){
    break
  }else{
    mu.lower <- mu.lower - mu0.iterate
  }
}

mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))
  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- closer.dat - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>

```

```

    mutate(xbars = xbars + mu.upper) # shifting back
# p-value
delta <- abs(mean(closer.dat) - mu.upper)
low <- mu.upper - delta # mirror
high<- mu.upper + delta # xbar
p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

if(p.val < 0.05){
  break
}else{
  mu.upper <- mu.upper + mu0.iterate
}
}

closer.rand.ci <- c(mu.lower, mu.upper)
# for difference data
R <- 1000
mu0.iterate <- 0.001
starting.point <- mean(difference.dat)
mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))
  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- difference.dat - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back
# p-value
delta <- abs(mean(difference.dat) - mu.lower)
low <- mu.lower - delta # mirror
high<- mu.lower + delta # xbar
p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- difference.dat - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)
    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back
# p-value
delta <- abs(mean(difference.dat) - mu.upper)
low <- mu.upper - delta # mirror
high<- mu.upper + delta
p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

if(p.val < 0.05){
  break
}else{
  mu.upper <- mu.upper + mu0.iterate
}
}

```



```

}
difference.rand.ci <- c(mu.lower, mu.upper)
#####

```

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the  $T$  statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples ( $s_r$ ), instead of the original sample ( $s$ )

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- (a) Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.  
We see that conducting the hypothesis test by using the standard deviation of the resample increases the Type I error rate significantly (from approximately 5% to 12%).

```

#####
R <- 10000
n <- 25
alpha <- 0.05
mu <- 0
type1_fixed <- numeric(R)
type1_resample <- numeric(R)
for (i in 1:R) {
  #x <- resamples.furtherft.stat
  x <- rnorm(n, mean = mu)
  s_fixed <- sd(x) # using original sd
  t_fixed <- mean(x) / (s_fixed / sqrt(n))
  resample <- sample(x,
                    size = n,
                    replace = TRUE)
  s_resample <- sd(resample) # using resample sd
  t_resample <- mean(resample) / (s_resample / sqrt(n))
  type1_fixed[i] <- (t_fixed < qt(0.05, n - 1))
  type1_resample[i] <- (t_resample < qt(0.05, n - 1))
}

mean(type1_fixed)

## [1] 0.0516

mean(type1_resample)

## [1] 0.1265

#####

```

- (b) Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?  
We see that we use the standard of the original data, the parameter of interest is captured 92.2% compared to 95.6% of the time when the resample standard deviation is used.

```

#####
R <- 1000
n <- 25
mu <- 0
coverage_fixed <- numeric(R)
coverage_resample <- numeric(R)
# function for CI using original sd
boot.stat.fixed <- function(d, i){
  s_fixed <- sd(x)
  mean(d[i]) / (s_fixed / sqrt(n))
}
# function for CI using resample sd
boot.stat.resample <- function(d, i) {
  x_i <- d[i]
  mean(x_i) / (sd(x_i) / sqrt(n))
}
for (i in 1:R) {
  x <- rnorm(n, mean = mu)

```

```

boots_fixed <- boot(data = x,
                   statistic = boot.stat.fixed,
                   R = 1000)
ci_fixed <- boot.ci(boots_fixed,
                   type = "bca")$bca[4:5]
coverage_fixed[i] <- (mu >= ci_fixed[1] & mu <= ci_fixed[2])
# boot ci using resample sd
boots_resample <- boot(data = x, statistic = boot.stat.resample, R = 1000)
ci_resample <- boot.ci(boots_resample, type = "bca")$bca[4:5]
coverage_resample[i] <- (mu >= ci_resample[1] & mu <= ci_resample[2])
}
mean(coverage_fixed)

## [1] 0.932

mean(coverage_resample)

## [1] 0.957

#####

```

## References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does  $n$  have to be for  $z$  and  $t$  intervals? *The American Statistician*, 54(2):121–128.